

Bapuji
Sr. Hadoop Developer
Phone: +1(224)-706-0020
Email: bapuji.dbj@gmail.com

PROFESSIONAL SUMMARY:

- Over **8+** years of experience including **4** years of Big Data Ecosystem related technologies with full project development, implementation and deployment.
- Strong Experience working with various Hadoop ecosystem components like, **Map Reduce, HDFS, Hive, Sqoop, Pig, Flume, and Oozie**.
- Strong Knowledge on **Architecture** of **Distributed** systems and **Parallel processing** frameworks.
- In-depth understanding of **MapReduce** Framework and **Spark** execution model.
- Worked extensively on fine-tuning long running **Spark** Applications to utilize better **parallelism** and **executor** memory for more caching.
- Strong experience working with both **batch** and **real-time** processing using **Spark** framework.
- Expertise in developing production ready **Spark** applications utilizing **Spark-Core, Data frames, Spark-SQL, Spark-ML** and **Spark-Streaming** API's.
- Hands on experience in installing, configuring and deploying **Hadoop** distributions in cloud environments (**Amazon Web Services**).
- Expertise in developing production ready **Spark** applications utilizing **Spark-Core, Data frames, Spark-SQL, Spark-ML** and **Spark-Streaming** API's.
- Worked on building real time data workflows using **Kafka, Spark streaming** and **HBase**.
- Worked extensively on **Hive** for building complex data analytical applications.
- Very good understanding of **Partitions, bucketing** concepts in **Hive** and designed both **Managed** and **External** tables in Hive to optimize performance.
- Used custom serDes like **Regex** SerDe, **JSON** SerDe, **CSV** SerDe etc., in hive to handle multiple formats of data.
- Having knowledge in **Apache Ambari** platform for securing, managing and monitoring Hadoop clusters.
- Experienced in Cluster coordination services through **zookeeper**.
- Strong experience using different columnar file formats like **Avro, RCFile, ORC** and **Parquet** formats.
- Worked with **Sqoop** to move (import/export) data from a relational database into Hadoop.
- Experience working with Hadoop clusters using **Cloudera, Amazon EMR** and **Hortonworks** distributions.
- Extensive experience in performing **ETL** on structured, semi-structured data using **Pig** Latin Scripts.
- Designed and implemented **Hive** and **Pig** UDF's using **Java** for evaluation, filtering, loading and storing of data.
- Experienced in job workflow scheduling and monitoring tools like **Oozie**.
- Well versed with **UNIX** and **Linux** command line and **shell** script.
- Adequate knowledge and working experience with **agile** methodology.

TECHNICAL SKILLS:

Languages	Java, Scala, SQL, PL/SQL, Pig Latin, Python, Hive QL
Web Technologies	JEE (JDBC, JSP, SERVLET, JSF, JSTL), AJAX, JavaScript
Big Data Systems	Hadoop, HDFS, MapReduce, YARN, Pig, Hive, Sqoop, Flume, Oozie, Impala, Spark, Apache Airflow, Kafka, Splunk, Cloudera CDH4, CDH5, Hortonworks, Hadoop EMR, Talend and Ranger...
RDBMS	Oracle 10g/11g, MySQL, SQL Server 2005/2008 R2, PostgreSQL, DB2, Teradata
NoSQL Databases	HBase, MongoDB, Cassandra
App/Web Servers	Apache Tomcat, WebLogic
SOA	Web services, SOAP, REST

Frameworks	Struts 2, Hibernate, Spring 3.x
Version Control	GIT, CVS, SVN
IDEs	Eclipse, Scala IDE, NetBeans, IntelliJ IDEA
Operating Systems	UNIX, Linux, Windows

EDUCATION:

Bachelor of Technology in Computer Science Engineering at JNTU, Kakinada, Andhra Pradesh, India.

WORK EXPERIENCE:

Cigna – Bloomfield, Connecticut
Role: Hadoop/Spark Developer

Jul'17 – Present

Responsibilities:

- Developed **Spark** applications using **Scala** utilizing **Data frames** and **Spark SQL** API for faster processing of data.
- Developed highly optimized **Spark** applications to perform various data **cleansing, validation, transformation** and **summarization** activities according to the requirement
- Data pipeline consists **Spark, Hive** and **Sqoop** and **custom built Input Adapters** to ingest, transform and analyze operational data.
- Developed **Spark** jobs and **Hive** Jobs to summarize and transform data.
- Used **Spark** for interactive queries, processing of **streaming** data and integration with **NoSQL** database **HBase, Cassandra** for interactive access patterns.
- Involved in converting **Hive** queries into **Spark** transformations using **Spark Data Frames** in **Scala**.
- Automated creation and termination of **AWS EMR** clusters using **AWS, java sdk**.
- Built **real time** data pipelines by developing **Kafka** producers and **spark streaming** applications for consuming.
- Ingested **syslog** messages to **Kafka**.
- Worked on **Apache Airflow** to schedule single and sometimes complex chains of tasks that depend on each other on regular intervals.
- Handled importing data from relational databases into **HDFS** using **Sqoop** and performing transformations using **Hive** and **Spark**.
- Having knowledge in **Apache Ambari** platform for securing, managing and monitoring Hadoop clusters.
- Exported the processed data to the relational databases using **Sqoop**, to further visualize and generate reports for the BI team.
- Experienced in cluster coordination services through **Zookeeper**.
- Installed, tested and deployed monitoring solutions with **Splunk** services.
- Used **Hive** to analyze the **partitioned** and **bucketed** data and computed various metrics for reporting.
- Developed **Hive** scripts in **Hive QL** to de-normalize and aggregate the data.
- Scheduled and executed workflows in **Oozie** to run various jobs.
- Designing & creating ETL jobs through **Talend** to load huge volumes of data into Cassandra, Hadoop Ecosystem and relational databases.

Environment: Hadoop, Spark, Hive, Java, Scala, Maven, Impala, Oozie, Oracle, Ambari, GitHub, Tableau, Unix, Hortonworks, Apache Airflow Kafka, Zookeeper, Sqoop, Cassandra, Talend, Splunk, HBase.

Qualcomm -- San Diego, CA
Role: Hadoop/Spark Developer

Dec'16 – Jun'17

Responsibilities:

- Part of **Big Data Center of Excellence (CoE)**, responsible for designing and building enterprise data analytics platform.
- Worked with respective business units in understanding the scope of the analytics requirements.
- Performed core **ETL** transformations in **Spark**.
- Automated data pipelines which involve data **ingestion**, data **cleansing**, data **preparation** and data **analytics**.
- Created end to end **Spark** applications using **Scala** to perform various data **cleansing**, **validation**, **transformation** and **summarization** activities on **user behavioral** data.
- Developed end-to-end data pipeline using **FTP Adaptor**, **Spark**, **Hive** and **Impala**.
- Implemented **Spark** utilizing **Spark-SQL** heavily for faster development, and processing of data.
- Exploring with **Spark** for improving the performance and optimization of the existing jobs in Hadoop using **Spark-SQL**, Data Frame running in **Yarn** mode.
- Handled importing other enterprise data from different data sources into **HDFS** using **Sqoop** and performing transformations using **Hive**, **Map Reduce** and then loading data into **HBase** tables.
- Collecting and aggregating large amounts of log data using **Flume** and staging data in **HDFS** for further analysis
- Wrapper developed in **Python** for instantiating multithreaded application and running with other applications.
- Analyzed the data by performing Hive queries (**Hive QL**) and running Pig scripts (**Pig Latin**) to study customer behavior.
- **Data warehousing**, experience in design, development and testing, implementation and support of enterprise **data warehouse**.
- Used **Hive** to analyze the partitioned and bucketed data and compute various metrics for reporting.
- Created components like **Hive UDFs** for missing functionality in **HIVE** for analytics.
- Worked on various performance optimizations like using **distributed cache** for small datasets, **Partition**, **Bucketing** in **Hive** and **Map Side joins**.
- Created **Oozie** workflows and coordinators to automate data pipelines daily, weekly and monthly.
- Automated creation and termination of **AWS EMR** clusters using **AWS**, java sdk.

Environment: AWS EMR, Hadoop, Spark, Hive, Sqoop, HBase, UNIX, Talend, Pig, Linux, Java, Scala, Python, Ambari, Zookeeper.

Hortonworks

McKesson - Alpharetta, GA
Hadoop/Spark Developer

Dec'15 – Nov'16

Responsibilities:

- Developed multithreaded **Java** based Input adaptors for ingesting **click stream data** from external sources like **ftp server** and **S3** buckets on daily basis.
- Created various **spark** applications using **Scala** to perform various enrichment of these click stream data combined with enterprise data of the users.
- Implemented batch processing of jobs using **Spark Scala API**.
- Developed **Sqoop** scripts to import/export data from **Oracle** to **HDFS** and into **Hive** tables.
- Stored the data in **columnar** formats using **Hive**.
- Involved building and managing **NoSQL** Database models using **HBase**.
- Worked in **Spark** to read the data from **Hive** and write it to **Hbase**.
- Optimized the **Hive** tables using optimization techniques like **partitions** and **bucketing** to provide better performance with **Hive QL** queries.
- Worked with multiple file formats like **Avro**, **Sequence**, **Parquet** and **Orc**.
- Converted existing **MapReduce** programs to **Spark** Applications for handling semi structured data like **JSON** files, **Apache** Log files, and other custom log data.

- Loaded the final processed data to **HBase** tables to allow downstream application team to build rich and data driven applications.
- Worked with a team to improve the performance and optimization of the existing algorithms in Hadoop using **Spark, Spark -SQL**, Data Frame.
- Implemented business logic in **Hive** and written **UDF's** to process the data for analysis.
- Used **Oozie** to define a workflow to coordinate the execution of **Spark, Hive** and **Sqoop** jobs.
- Addressing the issues occurring due to the huge volume of data and transitions.
- Designed, documented operational problems by following standards and procedures using **JIRA**.

Environment: Java, Hadoop 2.1.0, Map Reduce2, Spark, Unix, Pig 0.12.0, Hive 0.13.0, Linux, Sqoop 1.4.2, Flume 1.3.1, Eclipse, AWS EC2, and Cloudera CDH 4.

American Home Shield - Memphis, TN

Dec'14 – Nov'15

Role: Hadoop Developer

Responsibilities:

- Migrated the needed data from **MySQL** into **HDFS** using **Sqoop** and importing various formats of flat files in to **HDFS**.
- Mainly worked on **Hive** queries to categorize data of different claims.
- Involved in loading data from **LINUX** file system to HDFS
- Written customized **Hive** UDFs in **Java** where the functionality is too complex.
- Implemented **Partitioning**, Dynamic Partitions, Buckets in **HIVE**.
- Designing and creating **Hive** external tables using shared meta-store instead of derby with partitioning, **dynamic partitioning** and **buckets**.
- Generate final reporting data using Tableau for testing by connecting to the corresponding **Hive tables** using **Hive ODBC** connector.
- Responsible to manage the test data coming from different sources
- Reviewing peer table creation in **Hive**, data loading and queries.
- Weekly meetings with technical collaborators and active participation in code review sessions with senior and junior developers.
- Monitored System health and logs and respond accordingly to any warning or failure conditions.
- Gained experience in managing and reviewing **Hadoop** log files.
- Involved in scheduling **Oozie** workflow engine to run multiple **Hive** and **pig** jobs
- Involved **unit testing**, interface testing, system testing and user acceptance testing of the workflow tool.
- Created and maintained Technical documentation for launching **Hadoop Clusters** and for executing **Hive** queries and **Pig** Scripts

Environment: Apache Hadoop, HDFS, Hive, Map Reduce, Core Java, Pig, Sqoop, Cloudera CDH4, Oracle, MySQL.

Protective Life - Edina, MN

Oct'13 - Nov'14

Role: Java Developer

Responsibilities:

- Implemented a Web based Application using Servlets, **JSP**, spring, **JDBC**, **XML**.
- Involved in writing Spring Configuration **XML** file that contains declarations and other dependent objects declarations.
- Used hibernate to connect to Database to create the **DAO** layer.
- Developed Application Framework using Model-View-Controller using the technology Spring.
- Used **HTML**, **XHTML**, **XML**, **XSLT**, **XPATH**, **JSP** and Tag Libraries to develop view pages
- Multilayer Applications construction using Open **JPA**, **HTML5**, **Spring MVC**.

- Annotated **Spring** Architecture (Spring Beans)
- Implemented **UNIX shell** scripts to migrate various data files to S&P ratings repository
- Implemented smooth pagination capability using **JSP** to remove existing pagination utility
- Worked on **Geo API** to provide geological access capability to S&P.com site.
- Involved in **Agile** process to streamline development process with iterative development.
- Code reviews and Managing the **CVS** Repository.
- Prepare builds for **DEV** and **UAT** environments.
- Participating in the regular team meetings sprint planning meetings, user story review meetings etc.
- Involved in preparing High & low level design docs with **UML** diagrams using Microsoft **VISIO** tool.

Environment: JDK 1.5, XML, HTML, XHTML, JSP, Spring, DAO, Oracle Express edition, Apache ANT, CVS, Junit, UNIX, Log4J, CSS Style Sheets, Apache Tomcat, J2EE, Maven 3

Accenture – Hyderabad, India

Oct'11– Sep'13

Role: Java Developer

Responsibilities:

- Involved in Requirements **analysis, design, and development and testing.**
- Involved in setting up the different roles & maintained **authentication** to the application.
- Designed, deployed and tested **Multi-tier application** using the Java technologies.
- Involved in front end development using **JSP, HTML & CSS.**
- Implemented the Application using **Servlets**
- Deployed the application on **Oracle** Web logic server
- Implemented **Multithreading** concepts in java classes to avoid deadlocking.
- Used **MySQL** database to store data and execute **SQL** queries on the backend.
- Prepared and Maintained **test environment.**
- Tested the application before going live to production.
- Documented and communicated **test result** to the team lead on daily basis.
- Involved in weekly meeting with team leads and manager to discuss the issues and status of the projects.

Environment: J2EE (Java, JSP, JDBC, Multi-Threading), HTML, Oracle Web logic server, Eclipse, MySQL, JUnit.

Golan Technologies – Hyderabad, India

Jun'09 - Sep'11

Role: Java Developer

Golan Technologies range from turnkey solutions to custom, client-driven solutions in a variety of product categories including website development and platform based applications, demand intelligence and business insight generation. Smart sites have the ability to provide a unified user experience and consistent messaging on websites across the globe, driving a favorable brand impression.

Responsibilities:

- Involved in the analysis, design, implementation, and testing of the project.
- Developed UI using **HTML, JavaScript, CSS** and **JSP** for interactive cross browser functionality and complex user interface.
- Implemented the end-to-end functionality of the client requirement during the development phase.
- Implemented the functionality of **mapping** entities to the database using **Hibernate.**
- Written **SQL queries** involved in the **JDBC** connection in accordance with the business logic.
- Performed various levels of **unit testing** for the entire application using the **test cases**, which included preparation of detail documentation for the results.
- Actively participated in client meetings and taking the inputs for the additional functionality.

- Involved in fixing bugs and **unit testing** with test cases using **JUnit**.

Environment: J2EE, Spring, Hibernate, JavaScript, CSS, Servlets, MySQL